

UCLA Computer Science 33 (Fall 2016)

Midterm 1, 100 minutes, 100 points, open book, open notes.

Put your answers on the exam, and put your name and student ID at the top of each page.

Name: _____ Student ID: _____

1	2 a+b	2 c	3 a	3 b	4	5 a+b	5 c	5 d	total
									1111
									1110 0001

1 (11 minutes). In a circular shift, bits are not discarded when they fall off the end of a word; instead, they are reintroduced on the other end. Write a 64-bit function 'long rcshi (long a, int b);' for GCC on the x86-64 that returns the result of circularly shifting A right by B bits, where you can assume $0 \leq B < 64$. For example, $\text{rcshi}(\text{0xbaddeadbeefadded}, 16) == \text{0xdddedbaddeadbeefa}$.

16 shift \rightarrow 4 shift

```

long rcshi (long a, long b) {
    long mask = (a & (-1 << b)); // give last b bits
    long temp = a >> b; // make room for shift
    long s = mask << (64 - b); // more gained bits into right
                                // replace s
    return (s | temp); // replace s
}

```

}

Name: Timothy Marson Student ID: 004 800 078 (10)

2. On the x86-64, the 'cqto' instruction sets rdx to zero or to -1 depending on whether rax's sign bit is 0 or 1, respectively, and the 'idivq X' instruction divides the 128-bit signed integer $((2^{**64})*(long)rdx + ((unsigned long)rax))$ by the 64-bit signed integer X and puts the signed quotient into rax and the signed remainder into rdx. 'idivq X' traps if X is zero, or if integer overflow occurs when attempting to fit the quotient into rdx.

With that in mind, consider the following C code

7.5

```
long aquo (long a, long b) { return a / b; }
long arem (long a, long b) { return a % b; }
long bquo (long a, long b) { return -a / -b; }
long brem (long a, long b) { return -a % -b; }
```

When compiled for x86-64 by 'gcc -O2 -S', the compiler translates this source code into the following four functions, where the order and names of the functions have been changed. (Notice that A and B are identical.)

A:

```
aquo    movq    %rdi, %rax
        cqto    → exenu rdx
        idivq   %rsi
        ret
```

B:

```
aquo    movq    %rdi, %rax
        cqto
        idivq   %rsi
        ret
```

C:

```
arem    movq    %rdi, %rax
        negq    %rsi algae
        negq    %rax negale
        cqto    → exenu rdx
        idivq   %rsi
        movq    %rdx, %rax
        ret
```

D:

```
brem    movq    %rdi, %rax
        cqto
        idivq   %rsi
        movq    %rdx, %rax
        ret
```

2a (8 minutes): Label each machine-code functions (A, B, C, D) with the C-language function (aquo, arem, bquo, brem) or functions that it corresponds to.

8

2b (5 minutes): Explain why two of the C-language functions generate exactly the same machine code, even though mathematically they are different functions. Why isn't this a bug in the C compiler?

2

a/b and -a/-b will have the same machine code as long as a signed number, and the dividing is signed division.

Name: Timothy Mergano Student ID: 004 800 078

2c (10 minutes): Suppose we also use '-fwrapv', i.e., we compile with 'gcc -O2 -S -fwrapv'. Which part of the machine code for `aquo`, `arem`, `bquo`, and `brem` would you expect to change, and why? Give an example call showing why the given machine code would be incorrect if '-fwrapv' had generated it.

I would expect to change the two ~~leg~~ instructions
in the machine code for `brem`. (Think)
If `a` or `s` is the width, signed long, then
take the negative of its return. (Think)

7

Name: Timothy Marjoro

Student ID: 004 800 078

3. In this problem, your answers must use only the integer operations \wedge , $\&$, $|$, \sim , $!$, $==$, $!=$, $<$, $>$, $<=$, $>=$, $<<$, and $>>$ along with any integer constants that you find useful. Your answers should contain only straight-line code, i.e., no conditional expressions (?), conditional statements, or loops. You may assume that your code is compiled with -fwrapv. Minimize the number of operations that you use in your answers.

3a (10 minutes). Define a "strong integer" to be an integer whose binary representation contains two adjacent 1 bits, and a "weak integer" to be an integer that is not strong. Write a C function 'bool is_strong (long a);' that returns 1 if A is a strong integer, and 0 otherwise.

bool is_strong (long a) {
 int mask1 = 0x333333333333;
 int mask2 = 0x666666666666;
 int mask3 = 0x999999999999;
 int mask4 = 0xB BBB BBB BBB;
 int mask5 = 0xCCCCCCCCCCCC;
 int mask6 = 0xDDDDDDDDDDDD;

 long check1 = (a & mask1) | (a & (mask1 << 31));
 long check2 = (a & mask2) | (a & (mask2 << 31));
 long check3 = (a & mask3) | (a & (mask3 << 31));
 long check4 = (a & mask4) | (a & (mask4 << 31));
 long check5 = (a & mask5) | (a & (mask5 << 31));
 long check6 = (a & mask6) | (a & (mask6 << 31));

 long final = (check1 | check2 | check3 | check4 |
 check5 | check6);

 return (final != 0);
}

Name: Timothy Morgan Student ID: 004 800 078

3b (12 minutes). Write a C function 'long weakadd (long a, long b);' that returns the sum of two weak integers A and B. If the result would not fit in 'long', yield the low-order 64 bits of the correct mathematical answer. Remember, your implementation is limited to the operators mentioned on the previous page; in particular, it cannot use '+' or '-'.

$$-a = \sim a + 1$$

long weakadd (long a, long b) {

return a \oplus b;

(0)

}

Name: Timothy Morgan Student ID: 004 800 078

4 (12 minutes). The programming language Fortran, introduced in 1957 and still widely used in big scientific applications, has an arithmetic IF that uses a three-way branch, in which the numbers are labels of statements to go to depending on whether the if-expression is negative, zero, or positive. For example, assuming all quantities are 32-bit integers, this Fortran code:

```
if (M + N) 10, 20, 30  
10 I = I + 2  
20 J = J + 4  
30 K = K - 5
```

acts like this C code, where each "..." stands for $2^{31} - 4$ case labels:

```
switch (M + N) {  
    case -1: case -2: case -3: ...  
        I = I + 2;  
    case 0:  
        J = J + 4;  
    case 1: case 2: case 3: ...  
        K = K - 5;  
}
```

in that it adds 2 to I if $M + N < 0$, adds 4 to J if $M + N \leq 0$, and always subtracts 5 from K. Show how the above code can be translated to x86-64 machine code that uses only one comparison instruction, as opposed to the two or more comparisons that one might normally expect. Assume that M is in %rdi, %N is in %rsi, that I, J and K are global variables residing in RAM, and that -fwrapv is being used.

```
res %rsi  
cmpl %rsi, %rdi  
jle .L1  
je .L2  
jmp .L3
```

.L1 addl \$2, 8(%rsp) // $I = I + 2$
.L2 addl \$4, 16(%rsp) // $J = J + 4$
.L3 subl \$5, 24(%rsp) // $K = K - 5$

(12)

Name: Timothy Margone Student ID: 004 800 078

5. Consider the following x86-64 C program, which uses a flexible array member:

```
#include <stddef.h>
#include <stdlib.h>
#include <string.h>
struct cs {
    int color;
    long len;
    char data[];
};

/* Set *P to be a newly allocated string that has the old *P's
   color, but has only the bytes in the old *P starting at OFFSET
   and continuing for LEN bytes.  Return the address of the newly
   allocated string's data. */
char *
substr (struct cs **ptr, long offset, long len) {
    struct cs *old = *ptr;
    int a = alignof (struct cs);
    struct cs *new
        = malloc (offsetof (struct cs, data) + len);
    new->color = old->color;
    new->len = len;
    *ptr = new;
    /* The standard function memcpy (A, B, C)
       copies C bytes from B to A and returns A. */
    return memcpy (new->data, old->data + offset, len);
}
```

(6)

5a (3 minutes). Give the sizes and offsets of each member of 'struct cs', and why they have the values they do.

color : offset 0 ⁴ bytes. color is an int.

2. len : offset 4 ⁸ bytes. len is a long.

data : offset ¹⁶ unknown, size unknown. The size of the array is 4 known.

5b (5 minutes). Explain why the offsetof call is needed, how it works, and what it returns.

The offset call is needed to find the correct amount of memory to assign for the struct. It works by

casting the nullptr to type CS* and following it to the header data. The address of it is then taken and casted to (size_t) (like this: (size_t)&(((CS*)0)->data))

It will return the offset of the data array in memory.

Name: Timothy Mayon Student ID: CD 4 800 078

5c (12 minutes). Compiling the 'substr' function on the x86-64 might yield the following code, except that four faults (errors in the machine code) have been deliberately introduced. These faults are labeled A, B, C and D below. Fix each of the faults by correcting the machine code.

substr:

p 216

pushq %r13
movq %rsi, %r13
pushq %r12
movq %rdi, %r12
pushq %rbp
pushq %rbx
movq %rdx, %rbx
subq \$8, %rsp
A: movl (%rdi), %ebp
leaq 23(%rdx), %rdi
andq \$-8, %rdi
B: jle malloc
movl 0(%rbp), %edx
leaq 16(%rbp,%r13), %rsi
movq %rbx, 8(%rax)
leaq 16(%rax), %rdi
movl %edx, (%rax)
movq %rbx, %rdx
C: movl %eax, (%r12)
addq \$8, %rsp
popq %rbx
popq %rbp
D: popq %r13
popq %r12
jmp memcpy

ptr → %rdi
offset → %rsi
len → %rdx
movq (%rdi), %rbp + 3
call malloc + 3
73 + len = 7h
-8 3 rdh

(9)

Name: Thierry Lengow

Student ID: 004 80007P

(6)

5d (12 minutes). For each fault in (5c), give an example of what can go wrong in a C program if all the other faults are fixed but that fault remains unfixed. Be as precise as you can in your example.

Program

- A) Won't compile as incorrect suffix was ~~truncates~~ It will work but
- +1 ✓ ~~Used pointers are R bytes and need the suffix q. Also the register %rbp needs to be changed to its 64-bit version %rbp~~
- B) malloc is a function and it called everytime subtr is run. If malloc has a chance to not call malloc and the right amount of storage might never be allocated.
- C)
- D) The stack follows a last in, first out structure. The wrong value would be popped into a register %r13.
- +2 for the caller function