

CS33: Intro Computer Organization
Midterm, Form: B39

Name: _____

ID: _____

Please wait until everyone has their exam to begin. We will let you know when to start. Good luck!

Problem	Score	Points Possible
1	12	18
2	8	8
3	9	12
4	20	20
5	12	15
6	4	0
7	15	17
8	21	10

101

Question 1. The bigger the better. (18, 3 pts each)

12

1. What is the largest number that can be represented by a 7 bit floating point number (say with the same rules as IEEE 754 floating point), with a 1 bit sign, 3 bit exponent, and 3 bit significand (bias=3)? 15
2. In C, what's the largest `int` plus one? 2^{31} ~~2^{31}~~ -2^{31}
3. Consider an n-bit signed number, what's the largest one? 2^{n-1} ~~2^{n-1}~~
4. In C, what's the smallest `unsigned int` minus one? $2^{32}-1$ ~~$2^{32}-1$~~
5. Which can represent the largest number in C, the largest `float` or the largest `signed long` or largest `unsigned int`? signed long ~~signed long~~ largest float
6. Which integer type in C is large enough to store a pointer without loss of precision? long

1) looks like

0 110 111

$$\text{this is } (-1)^0 \times 2^{e-\text{bias}} \times (1+f) \quad \text{bias} = 2^{3-1} - 1 = 4 - 1 = 3$$

$$= 2^{4-3} \times (1 + \frac{1}{2} + \frac{1}{4} + \frac{1}{8}) = 8 \times (1 \frac{7}{8})$$

$$= 8 \times (1 \frac{7}{8}) = 15$$

Question 2. Matchmaker (8 Pts, 1 pts each)

Pretend to be a compiler.

You are free to assign registers to variables however you choose. Assume x and y are of type int. Remember, the compiler(me) may have done some optimizations.

- | | |
|--|-------------------------|
| <u>g</u> y=x+y | (a) shl \$ 5 %edi |
| <u>a</u> x=x*32 (x in %edi) | (b) xorl %edi %edi |
| <u>f</u> x=x*5+3 (x in %edi) | (c) shr \$ 31 %edi |
| <u>c</u> x=(x < 0) ? -1 : 0 | (d) movl \$1 %eax |
| <u>d</u> x=1 (x now in %eax) | (e) imul %edi %edx |
| <u>h</u> x=x*3+5 (x in %edi) | (f) leaq 3(%edi,%edi,4) |
| <u>b</u> x=0 (xor on x^x cancels itself) | (g) addl %edi %edi |
| <u>e</u> x=x*y (x in %edx, y in %edi) | (h) leaq 5(%edi,%edi,2) |

Question 3. Unholy Union (9 pts)

```
#include <stdio.h>
#include <string.h>
```

```
void main(char** argv, int argc) {
    union U {
        char s[16];
        int i;
        char c;
    } u;

    strcpy(u.s, "evil_prof"); //Copy string to destination from source

    printf("%x\n", u.c);
    printf("%x\n", u.i);
}
```

1. What does this program print? (6 pts)

0x65
0x65 76 69 6C
 e v l p

2. To which addresses may this union be aligned? (3pts)




can be aligned to address that
 are a multiple of 16 (end with 000₂)

This aligns to 4.

Question 4. Deconstructed (20 pts, 5 Each)

```
#include <stdio.h>
```

```
typedef struct {
    char a;
    int b;
    char c;
    double d;
} X;
```




```
void main(char** argv, int argc) {
    X x[10];
    printf("%d\n", (int)sizeof(X));
    printf("%d\n", (int)sizeof(x));
}
```

1. What does this program print?

X is 24 bytes, so prints 24
x is 240 bytes, so prints 240.

2. Draw the memory layout of X, where your diagram indicates which byte offset each variable is located at, as well as any space allocated just for padding:




3. Write an assembly snippet that performs `x[10].c=0`. Assume that x is in register \$rdi.

```
movq $0xA, %rax # 10 into rax.
imul $0x18, %rax # multiply 10 by 24.
movq $0, (%rdi,%rax) # set x[10].c (offset of 0x8 because that's where c is).
```

4. Describe how you would reduce the memory consumption of x. How small can you make x?

* Arrange fields with largest data types first:

```
typedef struct {
    double d;
    int b;
    char a;
    char c;
}
```



The struct itself can be 14 bytes,
but is aligned by 16.

Thus, can make array 160 bytes.

Question 5. I can puzzle, (15 Pts, 2 pts each)

Answer these true false puzzles. Assume the following setup:

```
int x = foo();
int y = bar();
unsigned ux = x;
unsigned uy = y;
```

- ✓ F $x > 0 \ \&\& \ y > 0 \Rightarrow x + y > 0$ No, can get overflow to negatives.
- ✓ F $5 * ux > ux$ $ux = 1000$ $5 * ux = 40 \Rightarrow 1010000 \Rightarrow 1000$, not > 1000
- ✓ F $x < 100 \Rightarrow 10 * ux > ux$ $x = 0$.
- ✓ T $-x == \sim x + 1$
- 3 T $x \gg 2 == x / 4$

Question 6. ... and so can you! (Up to 4 pts Extra Credit)

1. Write a C Puzzle of the form above, give the solution, and explain why you think its cool.

$$(\sim(x \& y)) \& \sim((\sim x) \& (\sim y)) == 1 \quad +4$$

\Downarrow

$!(x == y)$

true

- $(\sim(x \& y)) \& \sim((\sim x) \& (\sim y))$ is a way to express $x \wedge y$, since $\sim(x \& y)$ creates a mask where all places x and y both have 1's goes to 0, and $\sim(\sim x \& \sim y)$ is another way to express $x | y$.

- if $x \wedge y$ is true (non zero), then the two cannot be equal.

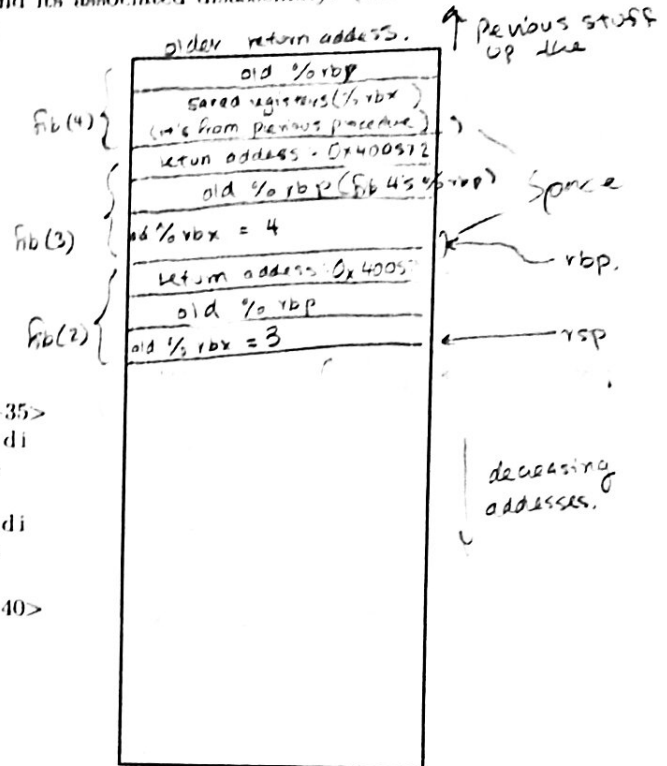
It's cool because it tests that you know to test for equivalence with $!(x \wedge y)$, and how \wedge works. Also, there are a lot of squiggles.

Question 7. Your fibs are stacking up (16 Pts)

Recall the fibonacci code that we discussed in class, and its associated disassembly: (the instruction addresses are omitted for simplicity, just the offsets remain)

```
int fib(int a) {
    if (a < 2) {
        return 1;
    }
    return fib(a-1) + fib(a-2);
}
```

```
fib: 0x40055d <+0>: push    %rbp
      0x40055e <+1>: push    %rbx
      0x40055f <+2>: sub     $0x8,%rsp
      0x400563 <+6>: mov     %edi,%ebx
      0x400565 <+8>: cmp     $0x1,%edi
      0x400568 <+11>: jle     0x400580 <fib+35>
      0x40056a <+13>: lea     -0x1(%rdi),%edi
      0x40056d <+16>: callq   0x40055d <fib>
heat 0x400572 <+21>: mov     %eax,%ebp
      0x400574 <+23>: lea     -0x2(%rbx),%edi
      0x400577 <+26>: callq   0x40055d <fib>
      0x40057c <+31>: add     %ebp,%eax
      0x40057e <+33>: jmp     0x400585 <fib+40>
      0x400580 <+35>: mov     $0x1,%eax
      0x400585 <+40>: add     $0x8,%rsp
      0x400589 <+44>: pop     %rbx
      0x40058a <+45>: pop     %rbp
      0x40058b <+46>: retq
```



1. This function calls itself recursively. Imagine in gdb we put a breakpoint on line 0x40056d, then call `fib(4)`. Furthermore we hit continue two more times in gdb, so that the stack frames of `fib(4)`, `fib(3)`, and `fib(2)` are all on the stack. Draw the contents of the stack in the box above, and be sure to indicate the stack pointer. Draw everything you know about the stack! If you know what the value is, write the value, otherwise indicate what it is. (10 pts)
2. On which line(s) (specify as offset from fib please!) is/are callee saved registers being saved? (1pt)
3. On which line(s) is/are callee saved registers being restored? (1pt)
4. On which line(s) is/are the input argument to fib being set? (1pt)
5. On which line(s) is/are the return value from fib being set (for the final time)? (1pt)
6. On which line(s) is/are the stack being allocated? (1pt)
7. On which line(s) is/are the stack being de-allocated? (1pt)

2. 0x40055e (0x40055d as well, if you consider %rbp callee-saved).
3. 0x400589 (0x40058a as well, if you consider %rbp callee-saved).
4. 0x40056a, 0x400574
5. 0x40057c
6. 0x400572, 0x40055d, 0x40055e.
7. 0x400589, 0x40058a.

Question 8. Oh Fuuuudge (10 pts)

You just finished your CS32 homework when all of a sudden you "rm -f my_homework.c". Thankfully, you didn't delete your binary file - phew. You forgot all the expressions in your source code, but you kind of remembered the overall structure. It's time to analyze the binary to fill out the remaining expressions.

```

<+0>: mov    $0x1, %r9d
<+6>: jmp     <func+54>
<+8>: movslq  %r9d, %rax
<+11>: mov     (%rdi, %rax, 4), %r8d
<+15>: lea     -0x1(%r9), %eax
<+19>: jmp     <func+28>
<+21>: mov     %edx, 0x4(%rdi, %rcx, 4)
<+25>: sub     $0x1, %eax
<+28>: test    %eax, %eax
<+30>: js      <func+43>
<+32>: movslq  %eax, %rcx
<+35>: mov     (%rdi, %rcx, 4), %edx
<+38>: cmp     %r8d, %edx
<+41>: jg      <func+21>
<+43>: cltq
<+45>: mov     %r8d, 0x4(%rdi, %rax, 4)
<+50>: add     $0x1, %r9d
<+54>: cmp     %esi, %r9d
<+57>: jl      <func+8>
<+59>: repz retq
  
```

Handwritten notes for assembly analysis:

- rax holds i* (next to line 8)
- r9d is i* (next to line 0)
- esi is n* (next to line 54)
- r8d is key* (next to line 11)
- rax holds j* (next to line 25)
- rcx holds j* (next to line 32)
- edx holds arr[j]* (next to line 35)
- more key into arr(j+4)* (next to line 45)
- test eax* (next to line 28)
- if reg eax jump 43* (next to line 30)

1. Fill in the code (2 Pts each .. Extra Credit Possible)

```

void func(int arr[], int n)
{
    int i, key, j;
    for (i = 0; i < n; i++)
    {
        key = arr[i];
        j = i - 1;
        while (j >= 0 && arr[j] > key)
        {
            arr[j+1] = arr[j];
            j = j - 1;
        }
        arr[j+1] = key;
    }
}
  
```

Handwritten notes for code completion:

- rdi is arr* (left margin)
- r9d is i* (left margin)
- esi is n* (left margin)
- r8d is key* (left margin)
- eax is j* (left margin)
- rdi is key* (right margin)
- i = 0* (right margin)
- l = 1* (right margin)
- arr[0] = key = 3* (right margin)
- j = 0* (right margin)
- arr[0] = 1* (right margin)
- arr[1] = 3* (right margin)
- l = 2* (right margin)
- arr[1] = key = 2* (right margin)
- j = 1* (right margin)
- arr[1] = 3* (right margin)
- arr[2] = arr[1]* (right margin)

2. What well-known algorithm is this? (2 Pts Extra Credit)

Bubble Sort.

+