

A Computer Science 33 (Fall 2016)

Term 1, 100 minutes, 100 points, open book, open notes.  
Turn in your answers on the exam, and put your name and

80-20

1	2 a+b	2 c	3 a	3 b	4	5 a+b	5 c	5 d	total
									-b = ~b+1

1 (11 minutes). In a circular shift, bits are not discarded when they fall off the end of a word; instead, they are reintroduced on the other end. Write a 64-bit function 'long rcshi (long a, int b);' for GCC on the x86-64 that returns the result of circularly shifting A right by B bits, where you can assume  $0 \leq B < 64$ . For example, `rcshi (0xbaddeadbeefadded, 16) == 0xdddedbaddeadbeefa.`



8

```
long rcshi (long a, int b){
    long moved = a >> b;
    long move = a << (64 - b);
    long mask = (1 << 63) >> b;
    return (mask & move) | (~mask & moved);
}
```

SigNS Unreadable

✓

On the x86-64, the 'cqto' instruction sets rdx to zero or to -1 depending on whether rax's sign bit is 0 or 1, respectively, and the 'idivq X' instruction divides the 128-bit signed integer  $((2^{64})*(long)rdx + ((unsigned\ long)rax))$  by the 64-bit signed integer X and puts the signed quotient into rax and the signed remainder into rdx. 'idivq X' traps if X is zero, or if integer overflow occurs when attempting to fit the quotient into  $\frac{rdx}{rax}$ .

With that in mind, consider the following C code

```
long aquo (long a, long b) { return a / b; }
long arem (long a, long b) { return a % b; }
long bquo (long a, long b) { return -a / -b; }
long brem (long a, long b) { return -a % -b; }
```

When compiled for x86-64 by 'gcc -O2 -S', the compiler translates this source code into the following four functions, where the order and names of the functions have been changed. (Notice that A and B are identical.)

A: aquo  
movq %rdi, %rax  
cqto  
idivq %rsi  
ret

B: bquo  
movq %rdi, %rax  
cqto  
idivq %rsi  
ret

C: brem  
movq %rdi, %rax  
negq %rsi  
negq %rax  
cqto  
idivq %rsi  
movq %rdx, %rax  
ret

D: arem  
movq %rdi, %rax  
cqto  
idivq %rsi  
movq %rdx, %rax  
ret

2a (8 minutes): Label each machine-code functions (A, B, C, D) with the C-language function (aquo, arem, bquo, brem) or functions that it corresponds to.

✓

A: aquo      B: bquo      C: brem

8

D: arem

2b (5 minutes): Explain why two of the C-language functions generate exactly the same machine code, even though mathematically they are different functions. Why isn't this a bug in the C compiler?

$$\frac{a}{b} = \frac{-a}{-b} \quad \text{b/c the negatives cancel, so}\quad 2$$

are equivalent

(10 minutes): Suppose we also use '-fwrapv', i.e., we compile with  
gcc -O2 -S -fwrapv'. Which part of the machine code for aquo, arem,  
bquo, and brem would you expect to change, and why? Give an example  
call showing why the given machine code would be incorrect if  
'-fwrapv' had generated it.

-fwrapv is a flag that allows wrapping around for overflow

3

In this problem, your answers must use only the integer operations , &, |, ~, !, ==, !=, <, >, <=, >=, <<, and >> along with any integer constants that you find useful. Your answers should contain only straight-line code, i.e., no conditional expressions (?::), conditional statements, or loops. You may assume that your code is compiled with fwrapv. Minimize the number of operations that you use in your answers.

1011  
01011001  
01001010  
01011100  
01000110  
0011

3a (10 minutes). Define a "strong integer" to be an integer whose binary representation contains two adjacent 1 bits, and a "weak integer" to be an integer that is not strong. Write a C function 'bool is\_strong (long a);' that returns 1 if A is a strong integer, and 0 otherwise.

```
bool is_strong (long a){
```

```
    long res = a & (a>>1);
```

```
    return !!res;
```

(7)

*cast to unsigned long.*

0000  
0001  
0011  
0010  
0111  
0101  
0110  
0100

84-25

(12 minutes). Write a C function 'long weakadd (long a, long b);' that returns the sum of two weak integers A and B. If the result would not fit in 'long', yield the low-order 64 bits of the correct mathematical answer. Remember, your implementation is limited to the operators mentioned on the previous page; in particular, it cannot use '+' or '-'.

long weakadd (long a, long b){

long o = (a & b) << 1;

long or = a | b;

long and = a & b;

long res = (o | or) & ~and;  
return res;

}

⑫

0010 0010  
0100

---

0101 0101  
0:1010 or:0101  
and:001  
res:101

---

0101 0001  
0:0010 and:0001  
or:0101  
res:0110

---

0101 0010 and:0000  
0:0000  
or:0101  
res:0111

---

1001 0101

and: 0001  
0: 0010  
or: 1101

res: 1110

-7 + 5

-9 + 4 + 2

84/85

(12 minutes). The programming language Fortran, introduced in 1957 and still widely used in big scientific applications, has an arithmetic IF that uses a three-way branch, in which the numbers are labels of statements to go to depending on whether the if-expression is negative, zero, or positive. For example, assuming all quantities are 32-bit integers, this Fortran code:

```
if (M + N) 10, 20, 30  
10 I = I + 2  
20 J = J + 4  
30 K = K - 5
```

acts like this C code, where each "..." stands for  $2^{31} - 4$  case labels:

```
switch (M + N) {  
    case -1: case -2: case -3: ...  
        I = I + 2;  
    case 0:  
        J = J + 4;  
    case 1: case 2: case 3: ...  
        K = K - 5;  
}
```

-1, does all 123  
0, does 23  
1, does 3

in that it adds 2 to I if  $M + N < 0$ , adds 4 to J if  $M + N \leq 0$ , and always subtracts 5 from K. Show how the above code can be translated to x86-64 machine code that uses only one comparison instruction, as opposed to the two or more comparisons that one might normally expect. Assume that M is in %rdi, %N is in %rsi, that I, J and K are global variables residing in RAM, and that -fwrapv is being used.

— because it always subtracts 5 from K, we only need to determine

if  $M + N = 0$ , then  $J = J + 4$ , else do  $J = J + 4$  and  $I = I + 2$

.L1  
movl %rdi, %rax

addl %rsi, %rax

test %rax, %rax

" $K = K - 5$ "  $\rightarrow$ ? machine instructions?

je .L2

" $I = I + 2$ "

ret

.L2

" $I = I + 2$ "

" $J = J + 4$ " and ~~K~~

⑧

Oct 29

Consider the following x86-64 C program, which uses a flexible array member:

```
#include <stddef.h>
#include <stdlib.h>
#include <string.h>
struct cs {
    int color;
    long len;
    char data[];
};

/* Set *P to be a newly allocated string that has the old *P's
   color, but has only the bytes in the old *P starting at OFFSET
   and continuing for LEN bytes.  Return the address of the newly
   allocated string's data. */
char *
substr (struct cs **ptr, long offset, long len) {
    struct cs *old = *ptr;
    int a = alignof (struct cs);
    struct cs *new
        = malloc (offsetof (struct cs, data) + len);
    new->color = old->color;
    new->len = len;
    *ptr = new;
    /* The standard function memcpy (A, B, C)
       copies C bytes from B to A and returns A. */
    return memcpy (new->data, old->data + offset, len);
}
```

4

5a (3 minutes). Give the sizes and offsets of each member of 'struct cs', and why they have the values they do.

int color size 4 bytes b/c need  $2^{32}$  to range int  
offset is 0 b/c is first  
long len size = 8 bytes b/c need  $2^{64}$  to range long  
offset is 4 b/c alignment is 8 and int already used 4  
data[ ]? 8 → 0-3

5b (5 minutes). Explain why the offsetof call is needed, how it works, and what it returns.

- 3
- needed to help allocate correct amount of memory. helps calculate how much memory the "int color; long len" require (includes alignment)
  - offsetof (struct cs, data) ~~returns~~ how much the header data needs  
len is how much the array needs
  - returns the ~~size~~ of how much memory it needs in bytes
    - it works by finding address of data and start of struct and ret the diff

7c-25

(12 minutes). For each fault in (5c), give an example of what can be wrong in a C program if all the other faults are fixed but that fault remains unfixed. Be as precise as you can in your example.

Q

Do not popping off stack correctly, so stack will break

X

(12 minutes). Compiling the 'substr' function on the x86-64 might yield the following code, except that four faults (errors in the machine code) have been deliberately introduced. These faults are labeled A, B, C and D below. Fix each of the faults by correcting the machine code.

substr:

```
pushq %r13  
movq %rsi, %r13  
pushq %r12  
movq %rdi, %r12  
pushq %rbp  
pushq %rbx  
movq %rdx, %rbx  
subq $8, %rsp  
A:  movl (%rdi), %ebp  
    leaq 23(%rdx), %rdi  
    andq $-8, %rdi  
B:  jle malloc  
    movl 0(%rbp), %edx  
    leaq 16(%rbp,%r13), %rsi  
    movq %rbx, 8(%rax)  
    leaq 16(%rax), %rdi  
    movl %edx, (%rax)  
    movq %rbx, %rdx  
C:  movl %eax, (%r12)  
    addq $8, %rsp  
    popq %rbx  
    popq %rbp  
D:  popq %r13 %or r12  
    popq %r12 %or r13  
    jmp memcpy
```

+3

③